

ORAL TECHNICAL PRESENTATION • PART III

From Data Sprawl to Decision Speed.

A governed path from the data the Agency already owns to decisions it can defend.

Enterprise Data and AI Enablement Services (EDAES) • GSA FEDSIM Task Order Request 26-EDAES-01 • Alliant 2 • on behalf of the Agency



Four presenters. Ninety minutes. Three topics in the order you asked for.

02



RUN OF SHOW • 90 MINUTES

Opening	8 min
Topic 1 • Technical approach	32 min
Topic 2 • Scenario response	22 min
Topic 3 • Management approach	24 min
Close	4 min

stopped precisely at 90, with a two-minute buffer

- Slides submitted with Part III; only slides presented are evaluated
- PWS reference top right; slide number lower right
- No transition effects; only Key Personnel present
- Handoffs by name; timer and compliance monitor on the team

LEAD PRESENTER
Dana Whitfield
Program Manager
Key Personnel • H.3.1

Rafael Okonkwo-Reyes
Chief Solutions Architect
Key Personnel • H.3.2

Priyanka Sethuraman
Data & AI Delivery Lead
Key Personnel • H.3.3

Marcus Lindqvist
Cyber Authorization Lead
Key Personnel • H.3.4

STRENGTH • FACTOR 4 • KEY PERSONNEL

The four presenters are the four Key Personnel who will run the task order; no proxy presenters.

The Agency does not need more pilots. It needs a governed path from data to decision.

03



THE AGENCY TODAY • TOR SECTION C

12

regional processing centers,
each with local data practices

~40M


benefit and compliance
transactions a year

300+

data sources across six legacy
case systems

9 → 0

AI pilots started since 2023;
none in production



*"I can get a prototype in a month. I cannot
get a decision I can defend in a year."*

Agency Director, industry day (illustrative)

Pilot counts are the illustrative TOR profile, not measured Agency data.


9 pilots built
notebooks, sample data,
one champion each

Gate 1 • Trusted data
governed, documented,
authorized access

**Gate 2 • Authorized
to operate**
inside an ATO boundary,
monitored

0 in production
the goal: used daily by
caseworkers

STRENGTH • FACTOR 1 • UNDERSTANDING The approach is organized around the two gates where the Agency's pilots actually stall, not around tools.




TOPIC 1 • TECHNICAL APPROACH • 32 MINUTES

Turn the two stall points into scheduled steps.

- **A governed data fabric with dataset passports**
Trusted data becomes reusable; Gate 1 becomes a step with an owner and a date.
- **A 90-day path from pilot to production**
Five gates, each decided by the Agency, with M-25-21 practices built in.
- **A security spine: inherit, automate, inspect**
Authorization measured in days, so Gate 2 keeps pace with delivery.

One governed data fabric. Three delivery lanes. One security spine under all of it.



STRENGTH • FACTOR 1 • TECHNICAL

A single reusable fabric feeds every product, so each new product inherits trusted data and authorization instead of rebuilding both.

Every PWS task has a method, a toolset, and a reason it fits this Agency.

TASK 1

Program management and transition

Single accountable PM, gated 60-day transition, integrated master schedule, planned-vs-actual cost reporting.

Why it fits: Twelve regions need one plan and one owner.

TASK 2

Enterprise data platform and governance

Metadata-first fabric: catalog, lineage, automated quality rules, attribute-based access; raw → governed → certified.

Why it fits: Gate 1 fails on trust and access, not storage.

TASK 6

Optional surge and innovation

Pre-qualified bench activated by modification; time-boxed experiments with exit criteria.

Why it fits: Adds capacity without diluting core delivery.



TASK 3

Analytics and AI enablement

90-day pilot-to-production path: feature store, model registry, evaluation harness, logged LLM gateway.

Why it fits: Nine pilots stalled at the same gates.

TASK 4

Security, authorization and zero trust

Inherit, automate, inspect: OSCAL control packages, policy-as-code, continuous monitoring, identity federation.

Why it fits: Delta packages replace whole-system re-authorization.

TASK 5

Adoption, training and change

Role-based learning paths, regional champions, in-product guidance, adoption telemetry.

Why it fits: A product no one uses does not shorten a queue.

One integrated offeror PWS · teal tasks remove the two historical stall points

STRENGTH · FACTOR 1 · TECHNICAL

Tasks 2 and 4 (teal) carry the methods that remove the two historical stall points; every other task supports them.

Every dataset earns a passport before anyone builds on it.



Pilot to production in 90 days, with a Government decision at every gate.

GATES G1-G5

The Agency product owner and, for high-impact uses, the Chief AI Officer's designee decide. Meridian Arc prepares the evidence and never self-approves.



1 • Intake

Mission owner states the decision to improve and the data involved

Exit: decision, owner and success metric named



2 • Value and risk triage

Value estimate; high-impact determination under OMB M-25-21; feasibility on certified data

Exit: go/no-go and required practices set



3 • Build on certified data

Feature store, model or workflow built only from certified datasets

Exit: reproducible build; lineage complete



4 • Evaluate

Accuracy, drift, bias and red-team tests; impact assessment for high-impact uses

Exit: results reviewed against thresholds



5 • Authorize

Delta control package into existing boundary; monitoring wired

Exit: authorizing official accepts



6 • Deploy and monitor

Workbench release, human oversight, feedback loop, appeals path

Exit: live with monitored metrics

STRENGTH • FACTOR 1 • TECHNICAL

A repeatable 90-day path with Government decisions at five gates converts stalled pilots into governed products.

Source: [OMB M-25-21](#), [OMB M-26-04 Unbiased AI Principles](#)



Inherit, automate, inspect: authorization that keeps pace with delivery.



inherited controls grow downward · new work shrinks upward



Each EDAES product

small delta package · target ≤ 15 business days to authorize



EDAES landing zone

control overlays, policy-as-code, OSCAL documentation, monitoring by default



Agency platform boundary

existing ATO; identity, network and logging controls inherited



FedRAMP-certified cloud services

infrastructure and platform controls inherited from the provider

THREE PRACTICES

Inherit

reuse provider and Agency controls; document once, reference everywhere

Automate

OSCAL packages and policy-as-code generated from the build pipeline, not written after it

Inspect

continuous monitoring views for the authorizing official; findings tracked to closure

ZERO TRUST RULES ON EVERY CONNECTION

Identity

Device

Network

Application

Data

Verify explicitly, least privilege, assume breach; all access to certified data logged to the Agency SOC.

STRENGTH · FACTOR 1 · TECHNICAL

Delta packages inside an existing boundary target authorization in days, not months, without lowering the authorizing official's bar.

Source: [FedRAMP: What is changing in 2026](#)



MERIDIAN ARC

Opening

Technical approach

Scenario response

Management approach

Close

9

Milestones the Agency can see and deliverables it can hold.



OPTION YEARS

OY1

Scale products to all regions; retire duplicate reports

OY2

Optimize cost and performance; automate certification

OY3

Extend to partner data; advanced decision support

OY4

Sustain, transfer knowledge, prepare recomplete

RECURRING DELIVERABLES

- **Transition-in plan and readiness report**
Day 5 • Day 60
- **Program management plan and IMS**
Day 30, then quarterly
- **Quality control plan (QASP-aligned)**
Day 30, annual update
- **Monthly status and cost report**
Monthly, 10th business day
- **Data governance playbook and passports**
Month 4, living
- **AI impact assessments and evaluation reports**
Per product, before Gate 5
- **Authorization delta packages**
Per product
- **Adoption and outcome dashboard**
Monthly, live

Option-year work is conditional on exercise and funding; milestones are planning targets confirmed at transition.

STRENGTH • FACTOR 1 • SCHEDULE

Six dated base-year milestones and eight recurring deliverables give the COR objective checkpoints from day 5 onward.

Metrics that change behavior, not just fill a monthly report.



≤ 5 days

Time to trusted data

business days from request to governed access for an approved role

GATE 1 • FROM CATALOG REQUEST LOGS



≤ 15 days

Authorization turnaround

complete delta package to authorizing decision

GATE 2 • FROM AUTHORIZATION TRACKER



+10% / qtr

Active product use

growth in weekly active users of released products until saturation

OUTCOME • FROM ADOPTION TELEMTRY



≥ 95

Certified quality score

weighted pass rate of automated quality rules on certified datasets

GATE 1 • FROM QUALITY ENGINE, WEEKLY



≤ 90 days

Product cycle time

intake to production for standard AI and analytics products

GATE 2 • FROM PRODUCT REGISTRY



≤ 2%

Defect escape rate

production defects found by users versus total defects found

OUTCOME • FROM SERVICE DESK AND QA LOGS

STRENGTH • FACTOR 1 • METRICS

Every metric has a target, a data source and a 10-day corrective-action rule; misses stay on the monthly report until closed.

What changes for a caseworker on a Tuesday morning.

TODAY

Three systems, one spreadsheet, first-in first-out

- Log in to three case systems; re-key into a spreadsheet
- Work the queue in arrival order; simple cases wait
- Ask a colleague which data is current
- Decide, then re-key the outcome into two systems
- No record of why one claim took 45 days, another 12



WITH EDAs

One workbench, a prioritized queue, a human decision

- One workbench fed by certified claims data
- Queue prioritized by completeness and urgency
- Each recommendation explains itself
- Caseworker decides; reason and decision logged together
- Leaders see the same certified numbers, by region

45 days

average time to a benefit-claim decision (illustrative baseline)

Decision speed

Topic 2 shows the path from 45 days to 15.

STRENGTH • FACTOR 1 • MISSION BENEFIT

The architecture is translated into a caseworker's day: one workbench, an explained queue, a logged human decision.

45 days to 15, in 120 days, with a human deciding every claim.

- **Sort and complete before you decide**
Thirty-seven of the 45 days are waiting, sorting and rework; a governed queue removes them.
- **Four gates, each decided by the Agency**
Scope and fairness criteria, evaluation results, authorization, outcome review.
- **Nothing new to authorize, nothing hidden**
Existing data, existing boundary, M-25-21 practices documented before scale



The scenario: 45 days to 15, in 120 days, without delaying a single planned authorization.

SCENARIO • ILLUSTRATIVE TOR

"The Agency Director has directed that the average benefit-claim decision time in the two largest regional centers fall from 45 days to 15 days within 120 days of task order award, using existing data sources only, without delaying any planned system authorization, and with documented compliance with OMB M-25-21 minimum risk management practices for high-impact AI. Describe how you would accomplish this."

Existing data only

no new integrations

No authorization slip

live inside current boundaries

M-25-21 practices


high-impact AI safeguards documented



Top: claims waiting to be sorted and completed • Bottom: claims decided by a caseworker

OUR READING • WHERE THE 45 DAYS GO TODAY

illustrative decomposition



Thirty-seven of forty-five days are waiting, sorting and rework. Those are the days a governed queue removes; the eight days of human review stay human.

WHAT SUCCESS LOOKS LIKE ON DAY 120

two regions on the governed queue • average decision time ≤ 15 days • every recommendation explained and logged • impact assessment, testing and appeals path documented • no authorization milestone moved

STRENGTH • FACTOR 1 • SCENARIO

The response attacks the 37 waiting days with a governed queue and leaves the 8 days of human judgment untouched.

Source: [OMB M-25-21](#)

A 120-day plan with four gates and a Government decision at each one.

DAYS 1-15

Discover and certify

- Certify claims intake data from the six case systems
- High-impact determination; impact assessment started
- Agree fairness metrics and thresholds

DAYS 16-45

Build the queue

- Completeness check and triage model on certified data
- Explanations written for caseworkers
- Pre-deployment testing: accuracy, bias, red team

DAYS 91-120

Scale and hand over

- All caseworkers in both regions on the workbench
- Outcome dashboard live for the Director
- Playbook, monitoring and appeals path handed over

GATE DECISIONS

G1 • Scope and fairness criteria confirmed

Decided by: Agency product owner

G2 • Evaluation results accepted

Decided by: Product owner + CAIO designee

G3 • Authorization decision and pilot metrics

Decided by: Authorizing official

G4 • Outcome review: ≤ 15 days

Decided by: Agency Director

DAYS 46-90

Pilot with humans in the loop

- Two regions; recommendations only, humans decide
- Delta authorization package inside existing boundary
- Weekly feedback and drift review

STRENGTH • FACTOR 1 • SCENARIO

Existing data sources only • existing platform boundary • humans make every decision • monitoring and appeals path live before scale.

How the governed queue works, and why a human still makes every decision.



Certified claims data

from six case systems; passport: owner, lineage, quality



Completeness and triage

check and triage model built only on certified data



Prioritized queue

each item explained in plain language



Caseworker decides

reviews the recommendation and reason, decides, records it; logged as one record

**Recommends, never decides.
Explains, always logs.**

MINIMUM PRACTICES, MAPPED

Pre-deployment testing

accuracy, bias and red-team results at G2

AI impact assessment

started day 1; complete before G3

Human oversight

caseworker decides; supervisors sample-review

Remedy and appeals

human re-review path live before scale

Ongoing monitoring

bias and drift weekly; pause rule if thresholds fail

Unbiased AI principles

truth-seeking, neutral outputs; no ideological weighting

Transparency


listed in the Agency AI inventory; plain-language explanations

STRENGTH • FACTOR 1 • SCENARIO

A caseworker decides every claim; recommendation, reason and decision are logged as one auditable record.

Source: [OMB M-25-21](#), [OMB M-26-04 Unbiased AI Principles](#)

Scenario risks, the triggers that reveal them, and what we need from the Agency.



GOVERNMENT

Product owner sets scope, fairness criteria and gate decisions • stewards certify datasets • AO decides the delta package • CAIO designee reviews high-impact determination • regional leads nominate champions

MERIDIAN ARC

Certify data, build the queue and explanations • run testing, assessment drafting and monitoring • prepare the delta package • train caseworkers, report weekly • hand over playbook, monitors and appeals



RISK AND TRIGGER

PRE-AGREED RESPONSE



Legacy data quality below certification threshold

Trigger: quality score < 90 in first 14 days

Certify only the fields the queue needs first; fix at source with stewards; publish gaps on the passport



Role and access approvals lag

Trigger: any role request open > 5 business days

Agree role catalog at G1; attribute-based policy pre-built; PM escalation within 48 hours



Fairness disparity beyond threshold

Trigger: disparity > agreed threshold at G2 or in weekly monitoring

Pause rule; root-cause analysis; re-evaluate before any scale; CAIO designee informed same day



Caseworker adoption stalls

Trigger: < 70% of pilot caseworkers active in week 3


Regional champions, in-product guidance, daily feedback, remove friction before adding features

STRENGTH • FACTOR 3 • RISK

Each scenario risk has an observable trigger and a pre-agreed response; fairness findings pause scale automatically.

Relevant work that puts this scenario in context.

ILLUSTRATIVE PROFILES • REPLACE WITH ACTUAL OFFEROR PAST PERFORMANCE



Benefits claims analytics modernization



Federal benefits program • illustrative

Consolidated claims data from legacy systems into one certified dataset; prioritized work queue with explanations; human-in-the-loop design with outcome monitoring.

Carries over: the queue pattern, explanation design and adjudicator training



Enterprise data fabric and certification model

Civilian agency CIO • illustrative

Catalog, lineage and quality rules across hundreds of datasets; steward-owned certification with published passports; attribute-based access tied to identity.

Carries over: passports, certification gates and steward roles



Continuous authorization pipeline

Civilian agency security office • illustrative

Machine-readable control documentation generated from the pipeline; delta packages inside an existing boundary; monitoring views for the AO.

Carries over: inherit, automate, inspect and the delta-package approach

STRENGTH • FACTOR 1 • SCENARIO Each profile names what carries over to this scenario, so evaluators can connect experience to the proposed approach.

One accountable chain. No surprises. Every month, in writing.

- **One point of accountability**

The Program Manager has authority over every person and partner on the task order.

- **Time-boxed escalation**

No issue waits more than two business days for an owner with authority to fix it.

- **Planned versus actual, monthly**

Cost, schedule, metrics, risks and staffing in one report the COR can act on.



One accountable chain from the AAS CO to every task lead.

GOVERNMENT

AAS Contracting Officer / COR
contract authority and performance oversight

Agency TPOC and product owners
priorities, gate decisions, stewardship


Authorizing official / CAIO designee
authorization and high-impact AI decisions

Program Manager • Dana Whitfield
single point of accountability; authority over all staff and partners • Key Personnel

Data Platform Lead (proposed function)
Task 2 • fabric, catalog, certification

Data & AI Delivery Lead • Priyanka Sethuraman
Task 3 • products, evaluation, gates • Key Personnel

Partner A (illustrative) • data engineering surge, inside Task 2



FORUMS

Daily task stand-ups

Weekly PM-COR status

Monthly program review with cost

Quarterly strategic review with the Director's office

Per gate product decisions

Chief Solutions Architect • Rafael Okonkwo-Reyes
deputy PM; technical baseline and PWS alignment • Key Personnel

Cyber Authorization Lead • Marcus Lindqvist
Task 4 • packages, monitoring, zero trust • Key Personnel

Adoption and Change Lead (proposed function)
Task 5 • training, champions, telemetry

Partner B • small business (illustrative) • training and change, owns Task 5

STRENGTH • FACTOR 3 • MANAGEMENT

A single Program Manager with authority over all staff and partners; the Government never has to ask which company owns a problem.

No issue waits more than two business days for an owner with authority to fix it.



Escalation ladder • issues climb only as far as authority requires

- SCHEDULED WITHIN 10 DAYS**
Level 4 • Joint review
Government and Meridian Arc leadership for scope, funding or authorization matters
- ≤ 5 BUSINESS DAYS**
Level 3 • Executive sponsor
corporate resources, partner decisions, Key Personnel actions; COR briefed same day
- ≤ 2 BUSINESS DAYS**
Level 2 • Program Manager
owns cross-task issues; reallocates staff or partner capacity; weekly status
- ≤ 1 BUSINESS DAY**
Level 1 • Task lead
owns, logs and works the issue; visible to the COR in the shared log

COMMUNICATION CADENCE		
Daily	Task stand-ups; shared issue and risk log	Task leads
Weekly	PM-COR status: progress, issues, next week	PM, COR, TPOC
Monthly	Status and cost report; program review; metrics	PM, COR, product owners
Quarterly	Strategic review; roadmap; adoption outcomes	Director's office, AAS CO
Per gate	Product gate decision packages	Product owner, CAIO designee, AO
Immediate	Problem notification within 24 hours of discovery	PM to COR and AAS CO


STRENGTH • FACTOR 3 • MANAGEMENT

Problem notification within 24 hours and a four-level ladder with time limits mean the Government always hears about a problem from us first.

Planned versus actual, by task, visible to the Government every month.

MONTHLY COST REPORT • PLANNED VS ACTUAL

illustrative format



One report, one dashboard.

Cost, schedule, performance, risk, staffing and look-ahead, every month, in writing.



Cost

planned vs actual hours and dollars by task and CLIN; funds status; invoice reconciliation




Schedule

IMS variance against milestones; critical path changes




Performance

the six metrics with trend and corrective actions




Risk and issues

register with triggers, owners and status changes



Staffing

fill rate, vacancies, pending substitutions



Look-ahead

next month's gates, deliverables and decisions needed

STRENGTH • FACTOR 3 • MANAGEMENT

Cost, schedule, metrics, risk and staffing arrive in one monthly report; variances beyond ±10% carry a recovery plan in the same report.

Staffed on day one. Resilient every day after.



Named backup for every role

cross-trained bench; backups attend stand-ups monthly so they know the work



Recruit before the vacancy


warm pipeline for the five hardest roles; offers ready within 10 business days of notice



Retain through the work

certification funding, rotation across tasks, clear paths from analyst to lead

ILLUSTRATIVE BASE-YEAR TEAM BY TASK (FTE SHARE)



≤ 30 days

target to fill any non-key vacancy; reported monthly as a coverage risk until closed (proposed)

KEY PERSONNEL SUBSTITUTION (H.3 PATTERN)

advance written notice • replacement meets or exceeds the qualification matrix • Government approval before the change • overlap funded by Meridian Arc • no Key Personnel changes in year one except for circumstances beyond our control

STRENGTH • FACTOR 4 • STAFFING

Named backups for every role, a 30-day fill target, and no year-one Key Personnel change without Government approval.

Sixty days to full performance, with no interruption the Agency can feel.

Full performance • day 60



Days 1–20 • Mobilize

Key Personnel on site day 1 • credentialing and access requests filed week 1 • transition-in plan day 5; joint kickoff • tooling and landing zone provisioned

Exit G1: access filed, plan accepted



Days 21–40 • Learn

Knowledge capture with incumbent and Agency staff • shadow current operations; document as-is • inventory nine pilots, their data and owners • risk register baselined; PMP draft

Exit G2: as-is documented, risks baselined



Days 41–60 • Lead

Assume responsibility task by task • PMP and QCP delivered by day 30; IMS baselined • first certified dataset published • readiness gate: staffing, access, tooling confirmed

Exit G3: readiness gate passed



HOW THE AGENCY KNOWS NOTHING BROKE

0

missed deliverables during transition

0

critical roles unfilled at day 60

≤ 5 days

oldest open access request

No change

to service levels or support hours

1

certified dataset before day 60

STRENGTH • FACTOR 4 • TRANSITION

Three 20-day phases with exit gates and five "nothing broke" measures the COR can verify at day 60.

Risks we already see, and the triggers that tell us when to act.



RISK • OWNER		TRIGGER	MITIGATION UNDERWAY
	Data access approvals slower than plan Data Platform Lead	Role request open > 5 days	Role catalog agreed at transition; policy pre-built
	Authorization package review time Cyber Authorization Lead	Delta package > 15 days without decision	Inherit and automate; weekly AO touchpoint; package quality checklist
	Legacy data quality below threshold Data Platform Lead	Certification score < 90	Field-level certification; fix at source with stewards
	Key staff attrition Program Manager	Any resignation notice	Named backups; warm pipeline; retention program
	Scope growth from regions Program Manager	Unplanned requests > 10% of capacity	Intake through product owner; change control; surge task if exercised

Register shared with the COR weekly; each pillar moves as its trigger is watched.

STRENGTH • FACTOR 3 • RISK

Each of five program risks has a trigger, a mitigation underway and a named owner; the register is shared weekly.

Quality is inspected in, measured monthly, and corrected within thirty days.



STRENGTH • FACTOR 3 • QUALITY

The QCP mirrors the QASP, so the Government and Meridian Arc measure the same standards from the same data at the same time.

Teaming partners are accountable the same way we are.



Workshare is an illustrative planning split; actual percentages follow the subcontracting plan.

Same standards, flowed down

subcontracts carry the QASP-aligned standards, the six metrics and the problem-notification rule verbatim

One team, one plan

partner staff sit inside task teams under Meridian Arc task leads; no separate partner workstreams

Monthly scorecard

metrics, staffing fill, deliverable acceptance and issue response per partner; quarterly executive review

Authority to act

PM may remove or replace partner staff; step-in rights if partner performance would put a Government metric at risk

MEANINGFUL SMALL BUSINESS WORK

Partner B owns Task 5 end to end: learning paths, champions network, adoption telemetry and quarterly feedback. That is a whole outcome with its own metric, not hours spread thin across tasks.



≈ 70%

Meridian Arc (prime) • PM, architecture, data platform, AI, security



≈ 18%

Partner A (illustrative) • data engineering surge, connectors



≈ 12%

Partner B • small business (illustrative) • Task 5 end to end

STRENGTH • FACTOR 3 • TEAMING

Partners sign the same standards, sit inside the same task teams and face the same scorecards; the PM holds replacement and step-in rights.

Why Meridian Arc: three commitments, mapped to how you will evaluate us.



One roof: three commitments, one accountable team.



Factor 1 • Technical approach and scenario

- Governed data fabric with dataset passports: Gate 1 becomes a scheduled step
- Pilot-to-production path with Government decisions at five gates
- Security spine: inherit, automate, inspect, so Gate 2 is measured in days
- 120-day scenario plan: 45 → 15 days with humans deciding and M-25-21 practices documented
- Six metrics with targets, sources and corrective actions



Factor 3 • Management approach

- One accountable chain from the AAS CO to every task lead
- No issue waits more than two business days for an owner
- Planned versus actual by task, visible monthly, with recovery plans
- Risk register with triggers and owners, shared weekly
- Partners held to the same standards, scorecards and step-in rights



Factor 4 • Key Personnel and staffing

- The four presenters today are the four Key Personnel on the task order
- Named backup for every role; recruit before the vacancy
- Sixty-day transition with a readiness gate and no felt interruption
- Substitution only with equal or better qualifications and Government approval
- Meaningful small business work: the whole adoption task

STRENGTH • ALL FACTORS

Every commitment above maps to a slide the evaluators have already seen and to a factor in Section M.9.

Ready for your clarification questions.

We have our Written Technical Proposal (Part III) on hand and will caucus only as long as needed to answer precisely. We will present nothing beyond the answers to the questions you pose.



Dana Whitfield
Program Manager



Rafael Okonkwo-Reyes
Chief Solutions Architect



Priyanka Sethuraman
Data & AI Delivery Lead



Marcus Lindqvist
Cyber Authorization Lead

Backup • Crosswalk: oral presentation topics to slides and evaluation factors.

Instruction	Required content	Slides	Evaluation
L.9.1(a)	Technical approach: rationale, methods and tools per PWS task; alignment to PWS	5, 6, 7, 8, 9	Factor 1
L.9.1(b)	Schedules of milestones and deliverables tied to the PWS	10	Factor 1
L.9.1(c)	Metrics with measurable contributions and how performance is managed	11	Factor 1
L.9.2(a)	Scenario: processes, methods, tools, relationships and qualifications	14, 15, 16	Factor 1
L.9.2(b)	Scenario: past and present work that puts the scenario in context	18	Factor 1
L.9.2(c)	Scenario: technical details, rationale and constructive plan	15, 16, 17	Factor 1
L.9.3(a)	Management: program management, process control, cost and status reporting, metrics	22, 26	Factor 3
L.9.3(b)	Management: governance, authority, communication, organization, problem resolution, teaming accountability	20, 21, 27	Factor 3
L.9.3(c)	Management: risk management and mitigation	17, 25	Factor 3
L.7.1 / L.7.2 / H.3	Staffing, Key Personnel and transition (also written volume)	2, 23, 24	Factor 4

Source: [GSA FEDSIM RFP 47QFCA24R0005 Sections L.8-L.9 and M.9 \(SAM.gov\)](#), [OCI: Developing orals for FEDSIM bids](#)

Backup • Illustrative task order baseline, research basis and content boundaries.

Task order (fictional)	GSA FEDSIM TOR 26-EDAES-01, Enterprise Data and AI Enablement Services, issued 2026 under Alliant 2 on behalf of "the Agency", a federal civilian benefits and compliance agency
Agency profile (illustrative)	12 regional centers • ~40M transactions a year • 300+ data sources across 6 legacy case systems • 9 AI pilots, none in production • 45-day average claim decision time
Period and structure	60-day transition-in • 12-month base • four 12-month options • offeror-developed PWS with Tasks 1–6
Key Personnel (fictional)	Program Manager • Chief Solutions Architect • Data & AI Delivery Lead • Cyber Authorization Lead
Oral presentation format	90-minute limit • Key Personnel present • Topics: Technical Approach, Scenario Response, Management Approach • slides submitted with Part III • PWS reference top right • slide number lower right • no transitions • Q&A with 30-minute caucus
Design decisions from orals research	first-90-seconds opening with discriminators • one repeated visual framework (the two gates on the decision path) • evaluator-language STRENGTH line on every content slide tied to a Section M factor • handoffs by name • isometric illustration of complex systems • consistent palette and icon family

Site G (47QFCA26R0010) was the only FEDSIM RFP publicly posted on SAM.gov in 2026; its Sections C/L/M are NDA-gated, so this case study uses a fictional task order modeled on the public FEDSIM template. Meridian Arc, its personnel, partners, experience and all figures are fictional and illustrative; no actual solicitation response, client, award, price or measured result is asserted.

REAL PUBLIC REFERENCES USED

[FEDSIM RFP 47QFCA24R0005 \(SWIB\)](#)
Sections L.7.7, L.8, L.9 and M.9

[FAR 15.102](#)
Oral presentations

[OMB M-25-21 and M-26-04](#)
Federal AI use; high-impact practices; unbiased AI principles

[FedRAMP 2026 changes](#)
Ongoing certification replaces continuous monitoring language

[OCI: FEDSIM orals, Parts 1 and 2](#)
Recurring pattern; slide-count and format discipline

[Lohfeld Consulting](#)
Evaluator-language strengths; scorecard-driven scripts

[Alpha Omega playbook](#)
First 90 seconds; one repeated framework; navigation cues

[APMP Western Chapter](#)
Consistent icons; isometric illustration of complex systems