

A useful agent starts with a human job.

The challenge

A delayed appliance delivery puts a customer's day on hold. Support has scattered order details and conflicting dates. How might we help the customer choose a reliable next step without repeating their story?

The customer's job

When my appliance delivery is delayed, I want to know my verified options and who owns the next step, so I can reorganize my day without chasing different teams.

Proposed starting point

Prototype one delivery-resolution agent with read-only tools and a human approval step. Compare a specialist agent team only where separate responsibilities improve the job.

How the workshop gets there

Discover the incident, define meaningful progress, map the job, design the agent responsibilities, prototype the conversation, test real scenarios and agree a bounded pilot.

The design logic

- Carry the story forward: D01 · Prototype one agent with a structured case brief; compare specialist handoffs.
- Match confidence to the evidence: D02 · Withhold an unverified promise and route conflicts to operations.
- Keep people in control: D03 · Require human approval before changing a booking or issuing a refund.

Make the next investment testable.

Scope and allocation

One delayed-delivery journey, sample order data, no live refunds or booking changes. A prototype study precedes a limited pilot.

Proposed: one facilitated workshop after research prework, then two weeks for prototype testing. A service owner, researcher, designer, agent builder and operations partner share the work.

Compare before committing

Use the same job and test set to compare one agent, a specialist team and a guided workflow. Separate roles only where that improves ownership, permissions or outcomes enough to justify the handoff overhead.

The evidence required next

Proposed formative study: 8 recent delayed-delivery customers and 4 frontline advisers. Include varied access needs. This is not a statistically representative sample.

Proposed gate: at least 6/8 customers complete critical tasks unaided; all 8 distinguish estimates from confirmed dates; at least 3/4 advisers can reuse the case. Every permission-boundary test must pass. Any unsupported promise or unauthorized action triggers redesign.

Decide through learning

- 30 days: Validate the job. Compare the prototypes. Interview recent customers, establish baselines and test one-agent and specialist designs with sample data.
- 60 days: Pilot one bounded workflow. Connect approved read-only tools for a limited group after the test gates pass. Monitor exceptions and human review.
- 90 days: Learn before expanding. Compare the pilot with the baseline and decide whether to revise, stop or widen the scope.

Meghan's contribution

Connect the customer story to the technical brief, facilitate cross-functional judgment, make the agent behavior visible and carry the evidence into an executive decision.

Illustrative AI agent workshop. Northstar Retail, research notes, scores and decisions are simulated. No live agents, customer systems or measured results.