

From a human job to a useful AI agent

Illustrative AI agent workshop. Northstar Retail, research notes, scores and decisions are simulated. No live agents, customer systems or measured results.

A human need. A clear build brief.

A delayed appliance delivery puts a customer's day on hold. Support has scattered order details and conflicting dates. How might we help the customer choose a reliable next step without repeating their story?

The proposed starting point

Prototype one delivery-resolution agent with read-only tools and a human approval step. Compare a specialist agent team only where separate responsibilities improve the job.

What participants leave with

- A customer job story and measures of useful progress.
- A task map with human judgment and agent responsibilities.
- A one-agent or specialist-team blueprint with explicit tools and permissions.
- A conversation prototype, scenario test set and pilot plan.

Scope

One delayed-delivery journey, sample order data, no live refunds or booking changes. A prototype study precedes a limited pilot.

Allocation

Proposed: one facilitated workshop after research prework, then two weeks for prototype testing. A service owner, researcher, designer, agent builder and operations partner share the work.

Prepare the room

Invitation

Join a facilitated AI agent design sprint: from the customer’s job to an agent blueprint, testable conversation and pilot plan. We will compare the simplest useful design with a coordinated agent team.

Pework

Before the session: gather consented recent customer incidents; anonymize relevant records; inventory available tools and data; bring service policies, a baseline and a sponsor decision. Without research access, treat all job stories as hypotheses.

People

Role	Contribution
Facilitator / service designer	Guide the job story, protect airtime and connect evidence to decisions.
Customer / research participant	Describe a recent incident and judge whether the prototype supports real progress.
Frontline adviser	Explain current workarounds and test proposed handoffs.
Agent builder / architect	Define context, tools, permissions and single-agent versus multi-agent choices.
Operations / data owner	Validate source reliability and take ownership of exceptions.
Service owner / sponsor	Agree outcomes, resource boundaries and the pilot decision.

Create space for good judgment

Room and materials

Use seven activity zones around one shared wall. Place customer evidence and the job story first, then the job map, agent contracts, prototype, test scorecard and pilot plan. Provide a shared digital equivalent.

- Customer-incident cards and JTBD canvas
- Job-map panels, sticky notes, pens and voting dots
- Agent role cards with input, output, tools and permissions
- Sample order records and service policies
- Conversation prototype and scenario test cards
- Outcome scorecard, decision log and pilot canvas

Working agreements

- Begin with the person and their circumstances.
- Separate evidence, interpretation and ideas.
- Write silently before discussion; make space for every voice.
- Preserve dissent and assign a test to the disputed assumption.
- Use consented, anonymized examples and minimum necessary data.

If the room gets stuck

Missing evidence becomes a named research task. Disagreement becomes a comparison using the same customer job and test cases. If time is short, reduce prototype scope; retain the test and decision gates.

Access and participation

Offer typed, spoken and remote ways to contribute. Obtain consent before recording or photographing. Use anonymized examples, identify who may see them and agree retention before collecting real research.

The day at a glance

09:00-14:50, with 65 minutes of breaks and lunch. The 285 minutes of activity assume completed research prework; a polished production agent is not the output of a single workshop.

Time	Activity	Output
09:00-09:35	Discover	Evidence of the job and its context
09:35-10:05	Define	A job story and success measures
10:05-10:15	Break	
10:15-10:55	Map	A task map with human and agent roles
10:55-11:45	Design	An agent blueprint and clear responsibilities
11:45-12:30	Lunch	
12:30-13:30	Prototype	A testable conversation and execution trace
13:30-13:40	Break	
13:40-14:25	Test	A test scorecard and prioritized changes
14:25-14:50	Pilot	A pilot brief, named owners and learning plan

01 / EMPATHIZE / 35 MIN

Start with the human story.

What we do

Find the progress a customer needs when a delivery goes wrong.

How we do it

Replay a recent incident. Listen for the trigger, workaround and emotional stakes.

Facilitator prompt

“I have taken time off work. I need a date I can actually plan around.”

Artifact to carry forward

Evidence of the job and its context

Interview prompts

- Tell me about the last delivery that did not go as planned.
- What happened just before you sought help?
- Walk me through the people and tools you tried, in order.
- What did you have to repeat or work out yourself?
- What made you trust or doubt an answer?
- What consequences did the delay have for your day and other people?
- What finally let you move forward? What remained unresolved?

02 / DEFINE / 30 MIN

Give the agent a job worth doing.

What we do

Turn the story into a clear job and a measurable outcome.

How we do it

Write “When... I want... so I can...” and agree what successful progress looks like.

Facilitator prompt

The goal is a dependable next step, with fewer repeated explanations.

Artifact to carry forward

A job story and success measures

Example job story

When my appliance delivery is delayed, I want to know my verified options and who owns the next step, so I can reorganize my day without chasing different teams.

Functional: Find a reliable next step and resolve the delivery exception without repeating context.

Emotional: Feel informed and in control when plans change.

Social: Coordinate confidently with the household and people covering work.

03 / DEFINE / 40 MIN

Find where AI can help.

What we do

Map the job, the information it needs and the decisions people own.

How we do it

Walk the current journey. Mark retrieval, judgment, handoffs and missing evidence.

Facilitator prompt

Retrieve the order. Check freshness. Explain options. Let a person approve a change.

Artifact to carry forward

A task map with human and agent roles

Task map

Recognize the problem: Capture the goal and ask only for missing essentials. No unnecessary personal data

Check the facts: Retrieve order and shipping records; inspect timestamps. No source, no invented answer

Compare options: Explain verified options and policy constraints. Label uncertainty and cite the source

Resolve & follow through: Draft a case summary and route to the right owner. No booking change or refund without approval

04 / IDEATE / 50 MIN

One agent. Or a team with purpose.

What we do

Design the smallest agent setup that can complete the job.

How we do it

Compare one agent with a specialist team. Define tools, context, handoffs and boundaries.

Facilitator prompt

Triage understands the request. Lookup checks the facts. Resolution proposes the next step.

Artifact to carry forward

An agent blueprint and clear responsibilities

Compare architectures

One agent + human approval: Simple to test and operate. Keep the job and tool set small enough to inspect.

Specialist agent team: Separates responsibilities, but adds handoffs, latency and coordination failure modes.

Guided workflow + adviser: A useful baseline when the process is predictable. Handles less ambiguity autonomously.

Agent contract

Name the job, input, output, permitted tools, handoff conditions and approval boundary. Every role needs an accountable owner.

05 / PROTOTYPE / 60 MIN

Turn the blueprint into an experience.

What we do

Make the agent's behavior visible before connecting real systems.

How we do it

Build a small conversation with sample tools, explicit uncertainty and a review step.

Facilitator prompt

See what the customer receives and how each agent reaches that response.

Artifact to carry forward

A testable conversation and execution trace

Build the smallest useful slice

- Use sample records before connecting customer systems.
- Make the customer conversation and tool trace visible together.
- Show source timestamps and explain uncertainty.
- Draft consequential actions; wait for authorized review.

Prototype instruction

Understand the delivery problem, retrieve only permitted records, explain verified options and draft the next step. Ask for missing context. Never invent a delivery commitment or execute a refund or booking change.

06 / TEST / 45 MIN

Prove the job gets done.

What we do

Test useful progress, including what happens when the data is wrong.

How we do it

Role-play realistic tasks. Observe comprehension, tool choices, recovery and handoffs.

Facilitator prompt

Try stale data, conflicting dates and a request outside the agent's permissions.

Artifact to carry forward

A test scorecard and prioritized changes

Observation plan

Observe unaided use. Capture task success, misunderstandings, repeated questions, tool traces and reviewer effort. Use consented, anonymized records.

Test first

- Clear source: Explain the unconfirmed status, cite the source and offer an adviser review.
- Stale data: Withhold the promise, show the stale timestamp and request verification.
- Conflicting dates: Surface the conflict and route it to operations; do not silently pick a date.
- Missing order: Ask for a permitted lookup detail or offer secure human support.
- Approval required: Draft the request and wait for authorized human review.

07 / LEARN & DELIVER / 25 MIN

Make the next step accountable.

What we do

Agree who will build, measure and improve the first bounded pilot.

How we do it

Review the evidence, name an owner and define advance, revise and stop gates.

Facilitator prompt

Begin with read-only access. Expand only when the evidence supports the next permission.

Artifact to carry forward

A pilot brief, named owners and learning plan

Decision readback

Name the chosen scope, its evidence, the dissent, the build owner and the next gate. A simulation can illustrate the plan but cannot validate a real agent.

Pilot boundary

One delayed-delivery journey, sample order data, no live refunds or booking changes. A prototype study precedes a limited pilot.

Source cards and traceability

Every note below is simulated. In a real workshop, replace each with a consented observation or clearly labeled assumption.

N01 / Evidence

I gave my order details to the store, then had to repeat everything to support.

Customer A | Simulated workshop input 1 | I01

N02 / Evidence

I cannot see which next step the previous adviser promised.

Frontline adviser | Simulated workshop input 2 | I01

N03 / Evidence

The app said Tuesday. Support said Thursday. I still did not know which was confirmed.

Customer B | Simulated workshop input 3 | I02

N04 / Evidence

A shipping scan is an estimate; it is not an approved delivery appointment.

Operations partner | Simulated workshop input 4 | I02

N05 / Evidence

Tell me the options, but do not change my order without asking me.

Customer C | Simulated workshop input 5 | I03

N06 / Evidence

I want to solve it myself if the answer is clear. Otherwise, give a person the context.

Customer D | Simulated workshop input 6 | I03

N07 / Interpretation

A compact, editable case brief could prevent repeated explanations.

Service designer | Simulated workshop input 7 | I01

N08 / Idea

Pass order ID, the customer's goal and verified facts between agent roles.

Agent builder | Simulated workshop input 8 | I01

N09 / Risk

A stale lookup could make a fluent answer sound more certain than the source.

Researcher | Simulated workshop input 9 | I02

N10 / Opportunity

Show the source and timestamp; escalate conflicting dates instead of choosing one.

Operations partner | Simulated workshop input 10 | I02

N11 / Open question

Who can authorize a booking change, and how is that approval recorded?

Service owner | Simulated workshop input 11 | I03

N12 / Idea

Human approval must be explicit before any consequential order change.

Service designer | Simulated workshop input 12 | I03

From evidence to design

I01 / Carry the story forward

Repeated explanations suggest a context gap. Test whether a shared case brief reduces rework without carrying unnecessary personal data.

Notes: N01, N02, N07, N08

O01 · Share the customer's goal and verified facts across permitted roles.

D01 · Prototype one agent with a structured case brief; compare specialist handoffs.

I02 / Match confidence to the evidence

Conflicting dates suggest that speed alone cannot establish trust. The response needs source freshness and a clear confirmation owner.

Notes: N03, N04, N09, N10

O02 · Retrieve the source, timestamp and uncertainty alongside the answer.

D02 · Withhold an unverified promise and route conflicts to operations.

I03 / Keep people in control

Some customers want self-service while others need help. Both need a visible boundary between an explanation and an authorized action.

Notes: N05, N06, N11, N12

O03 · Offer a reviewable next step and a context-preserving human handoff.

D03 · Require human approval before changing a booking or issuing a refund.

Agent responsibility cards

Triage

Understand the customer's goal and gather the minimum context.

Input: Request + order ID

Tools: Case-intake form

Output: Structured case brief

Boundary: Ask for missing details; do not infer identity or order ownership.

Lookup

Check order, shipping status and source freshness.

Input: Verified order reference

Tools: Read-only order lookup + shipment lookup

Output: Facts + sources + timestamps

Boundary: Flag stale, missing or conflicting records.

Resolution

Explain options and propose the next step.

Input: Case brief + verified facts

Tools: Approved policy search + draft-case tool

Output: Sourced response + proposed handoff

Boundary: Cannot confirm a new booking, refund or account change.

Human review

Resolve exceptions and authorize consequential changes.

Input: Customer choice + trace + proposed action

Tools: Approved service workflow

Output: Authorized action or request for more evidence

Boundary: Named adviser owns the decision and follow-through.

Compare options with the same rubric

Scores and weights are illustrative judgments. Weighted score = sum(score x weight) / sum(weights). Higher is better; changing the weights explores sensitivity.

Job completion (35%): 1 = partial answer; 5 = clear progress on the core job

Trust & control (25%): 1 = unclear evidence or authority; 5 = visible sources and approval

Ease of building (20%): 1 = complex dependencies; 5 = small, feasible first slice

Ease of operating (10%): 1 = hard to monitor; 5 = easy to inspect and support

Adaptability (10%): 1 = rigid; 5 = responsibilities can evolve independently

Design	Job completion	Trust & control	Ease of building	Ease of operating	Adaptability
One agent + human approval	4	4	5	5	3
Specialist agent team	5	4	2	2	5
Guided workflow + adviser	3	5	5	5	2

Test with people and the real implementation

A sourced, context-preserving agent helps a customer choose a reliable next step without surrendering control.

Participants

Proposed formative study: 8 recent delayed-delivery customers and 4 frontline advisers. Include varied access needs. This is not a statistically representative sample.

Observation

Observe unaided use. Capture task success, misunderstandings, repeated questions, tool traces and reviewer effort. Use consented, anonymized records.

Thresholds

Proposed gate: at least 6/8 customers complete critical tasks unaided; all 8 distinguish estimates from confirmed dates; at least 3/4 advisers can reuse the case. Every permission-boundary test must pass. Any unsupported promise or unauthorized action triggers redesign.

Next

Passing supports a limited pilot decision. Failing means revise the prompts, workflow, interface or underlying data, then test again. The on-site simulation is not evidence of live agent quality.

Decision log

D01 / Start with one agent and a structured case brief.

Evidence: I01 / N01, N02, N07, N08. Compare specialist handoffs with the same tasks.

Dissent: The architecture lead prefers separate agents for each responsibility.

Owner: Agent product owner

Open question: Does role separation improve outcomes enough to justify added coordination?

Next: Compare both prototypes on the same test set before choosing a pilot design.

D02 / Keep order tools read-only and expose uncertainty.

Evidence: I02 / N03, N04, N09, N10. Fluency cannot resolve conflicting sources.

Dissent: Operations asks whether visible uncertainty could increase contacts.

Owner: Operations + engineering leads

Open question: Which system can confirm a delivery slot, and how fresh must it be?

Next: Validate source ownership, freshness checks and the conflict escalation path.

D03 / Require human approval for consequential changes.

Evidence: I03 / N05, N06, N11, N12. Customers retain control of proposed actions.

Dissent: The sponsor wants to test more automation sooner.

Owner: Service owner

Open question: Can reviewers understand the evidence and complete handoffs reliably?

Next: Test approval comprehension, permissions and reviewer workload before expansion.

A gated 30 / 60 / 90 plan

30 days / Validate the job. Compare the prototypes.

Interview recent customers, establish baselines and test one-agent and specialist designs with sample data.

Owner: Research + service design + agent builder

Dependency: Participant access, approved sample data and named service owner

Measure: Task completion, repeated questions, uncertainty comprehension and trace quality

60 days / Pilot one bounded workflow.

Connect approved read-only tools for a limited group after the test gates pass. Monitor exceptions and human review.

Owner: Agent product owner + engineering + operations

Dependency: Source validation, access review, monitoring and a rollback path

Measure: Correct resolution, unsupported claims, escalation quality, latency and cost per resolved case

90 days / Learn before expanding.

Compare the pilot with the baseline and decide whether to revise, stop or widen the scope.

Owner: Sponsor + service owner

Dependency: Evidence of useful outcomes and sustainable support workload

Measure: Customer progress, resolution time, repeat contacts and reviewer effort

Method foundations

Original illustrative workshop combining JTBD discovery with design thinking and agent prototyping.

Jobs to Be Done theory · Christensen Institute / <https://www.christenseninstitute.org/theory/jobs-to-be-done/>

Rapid prototyping · IDEO.org Design Kit / <https://www.designkit.org/methods/rapid-prototyping.html>